



# Systems Software

Week 4: Orphans, Zombies and Deamons



# Overview

- ↗ Orphan Processes
- ↗ Zombie Processes
- ↗ Daemons
- ↗ Steps to create a daemon
- ↗ Required Reading

# Process Groups

- Every process running is part of a unique process group (**PGID**)
- When a process is created, it becomes a member of the group its parent is in.
- The first process member in the group sets the PGID to be equal to its PID.
- The first member in the group is referred to as the **Group Leader**

# Process Groups – System Calls

↗ Useful Process Group commands (see man for more details)

↗ `getpgrp()`

↗ `getpgid(0)`

↗ `getpgid(PID)`

↗ Shell Commands

↗ To list the PGID with the `ps` command use the `-j` flag

# Sessions and Session Groups

- All processes are grouped by sessions.
- These are linked to types of groups, eg. A user logs in, all processes the user has running are in its Session Group.
- When a user logs out, the kernel will terminate all processes in the user's session group.
- The session's ID is the same as the pid of the process that created the session through the `setsid()` system call.
- This session is referred to as the Session Leader.
- The `setsid()` system call takes no parameters and will return the new session id.

# Controlling Terminal

- Every session is associated with a terminal
- Processes in the session get Input and Output from the terminal
- A session is linked to a terminal (this is called the Controlling Terminal or Controlling tty)
- A terminal can be the controlling terminal for only one session at a time.
- An individual process disconnects from its controlling terminal when it calls `setsid` to become the leader of a new session.

# Controlling Terminal

- One of the attributes of a process is its controlling terminal.
- Child processes created with `fork` inherit the controlling terminal from their parent process.
- In this way, all the processes in a session inherit the controlling terminal from the session leader.
- A session leader that has control of a terminal is called the *controlling process* of that terminal.
- An individual process disconnects from its controlling terminal when it calls `setsid` to become the leader of a new session.

# File Descriptor

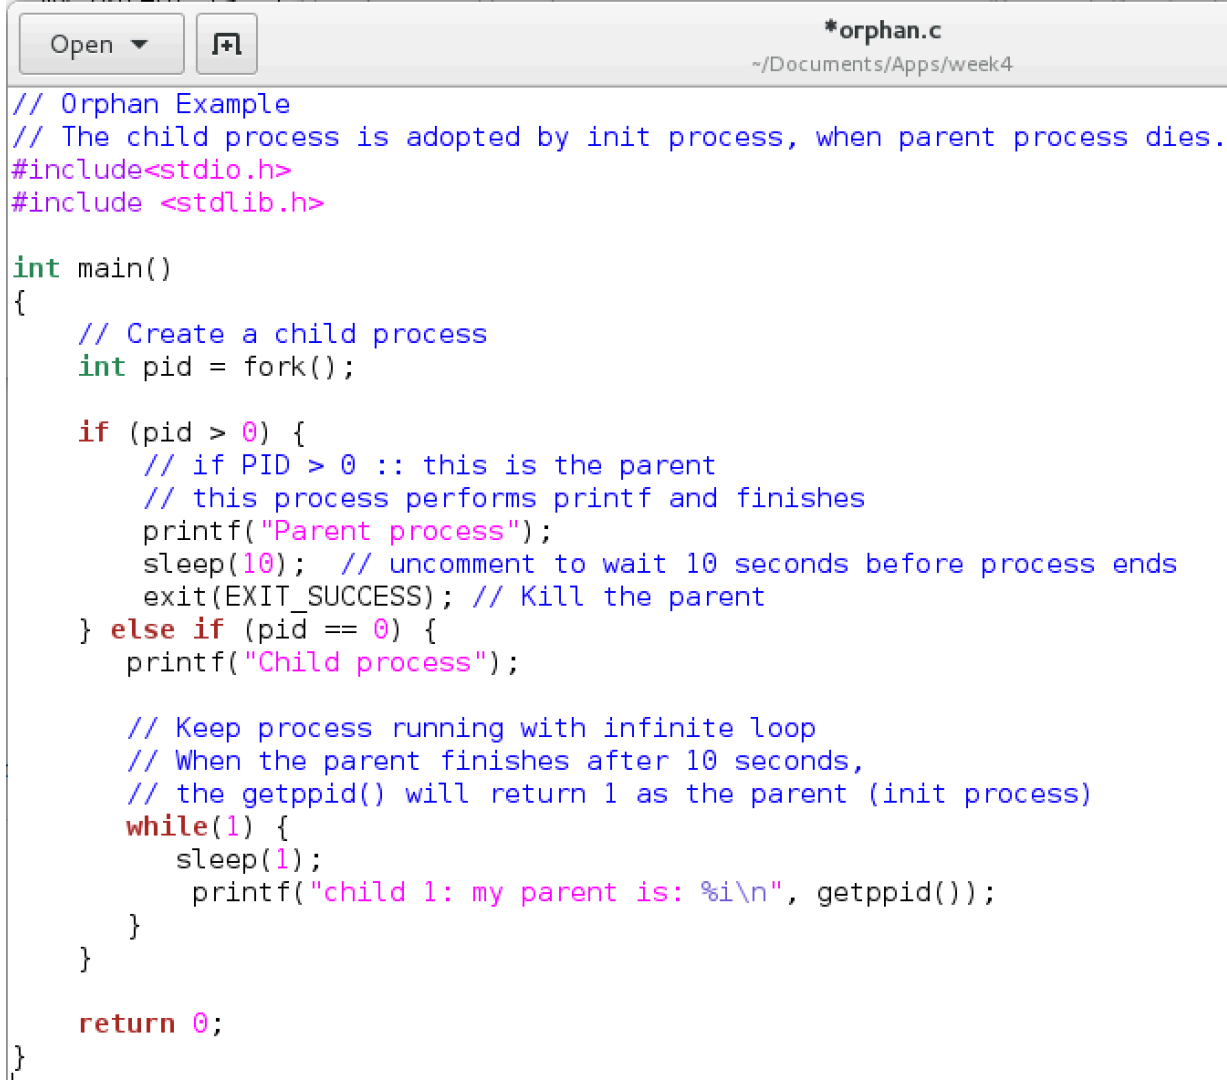
- Each time a file is opened, the OS creates a record to represent the file.
- This information is stored in the system kernel.
- An integer value is assigned to each record/entry.
- This integer value is the file descriptor. If a process has opened 5 files, there will be 5 file descriptors associated with the process.



# Orphan Processes

- An orphan process is a process whose parent has terminated.
- The orphan will default back to init as its parent (PPID of 1).
- The orphan will continue to run until it is killed:
  - Kill -9 PID
  - Where PID is the process ID of the orphan

# Example



```
// Orphan Example
// The child process is adopted by init process, when parent process dies.
#include<stdio.h>
#include <stdlib.h>

int main()
{
    // Create a child process
    int pid = fork();

    if (pid > 0) {
        // if PID > 0 :: this is the parent
        // this process performs printf and finishes
        printf("Parent process");
        sleep(10); // uncomment to wait 10 seconds before process ends
        exit(EXIT_SUCCESS); // Kill the parent
    } else if (pid == 0) {
        printf("Child process");

        // Keep process running with infinite loop
        // When the parent finishes after 10 seconds,
        // the getppid() will return 1 as the parent (init process)
        while(1) {
            sleep(1);
            printf("child 1: my parent is: %i\n", getppid());
        }
    }

    return 0;
}
```

# Run Orphan

jmccarthy@debianJMC2017: ~/Documents

File Edit View Search Terminal Help

\$. /orphan

jmccarthy@debianJMC2017: ~/Documents/Apps/wee

File Edit View Search Terminal Help

\$. /orphan

Child process child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

jmccarthy@debianJMC2017: ~/Documents/Apps/w

File Edit View Search Terminal Help

child 1: my parent is: 5308

child 1: my parent is: 5308

child 1: my parent is: 5308

Parent process\$ child 1: my parent is: 1

child 1: my parent is: 1

child 1: my parent is: 1

child 1: my parent is: 1

child 1: my parent is: 1

child 1: my parent is: 1

child 1: my parent is: 1

child 1: my parent is: 1

# Run Orphan

```

jmccarthy@debianJMC2017: ~/Documents/Apps/week4
File Edit View Search Terminal Help
PPID  PID  PGID  SID TTY      TPGID STAT  UID   TIME COMMAND

```

```

jmccarthy@debianJMC2017: ~/Documents/Apps/week4
File Edit View Search Terminal Help
1398  1499  1499  1499 pts/0      5342 Ss    1000   0:00 bash
1    1536  1046  1046 ?         -1  Sl    1000   0:00 /usr/lib/dconf/dconf-se
1145  2178  968   968 ?         -1  Sl    1000   7:17 firefox-esr
1398  2458  2458  2458 pts/1      5346 Ss    1000   0:00 bash
1    3996  1046  1046 ?         -1  Sl    1000   0:15 /usr/bin/gedit --gappli
1398  4122  4122  4122 pts/2      4122 Ss+   1000   0:00 bash
1499  5342  5342  1499 pts/0      5342 S+    1000   0:00 ./orphan
5342  5343  5342  1499 pts/0      5342 S+    1000   0:00 ./orphan
2458  5346  5346  2458 pts/1      5346 R+    1000   0:00 ps -xj
jmccarthy@debianJMC2017:~/Documents/Apps/week4$

```

Parent terminates  
after 10 seconds

Note the PPID of  
the orphan  
process has  
changed to 1

```

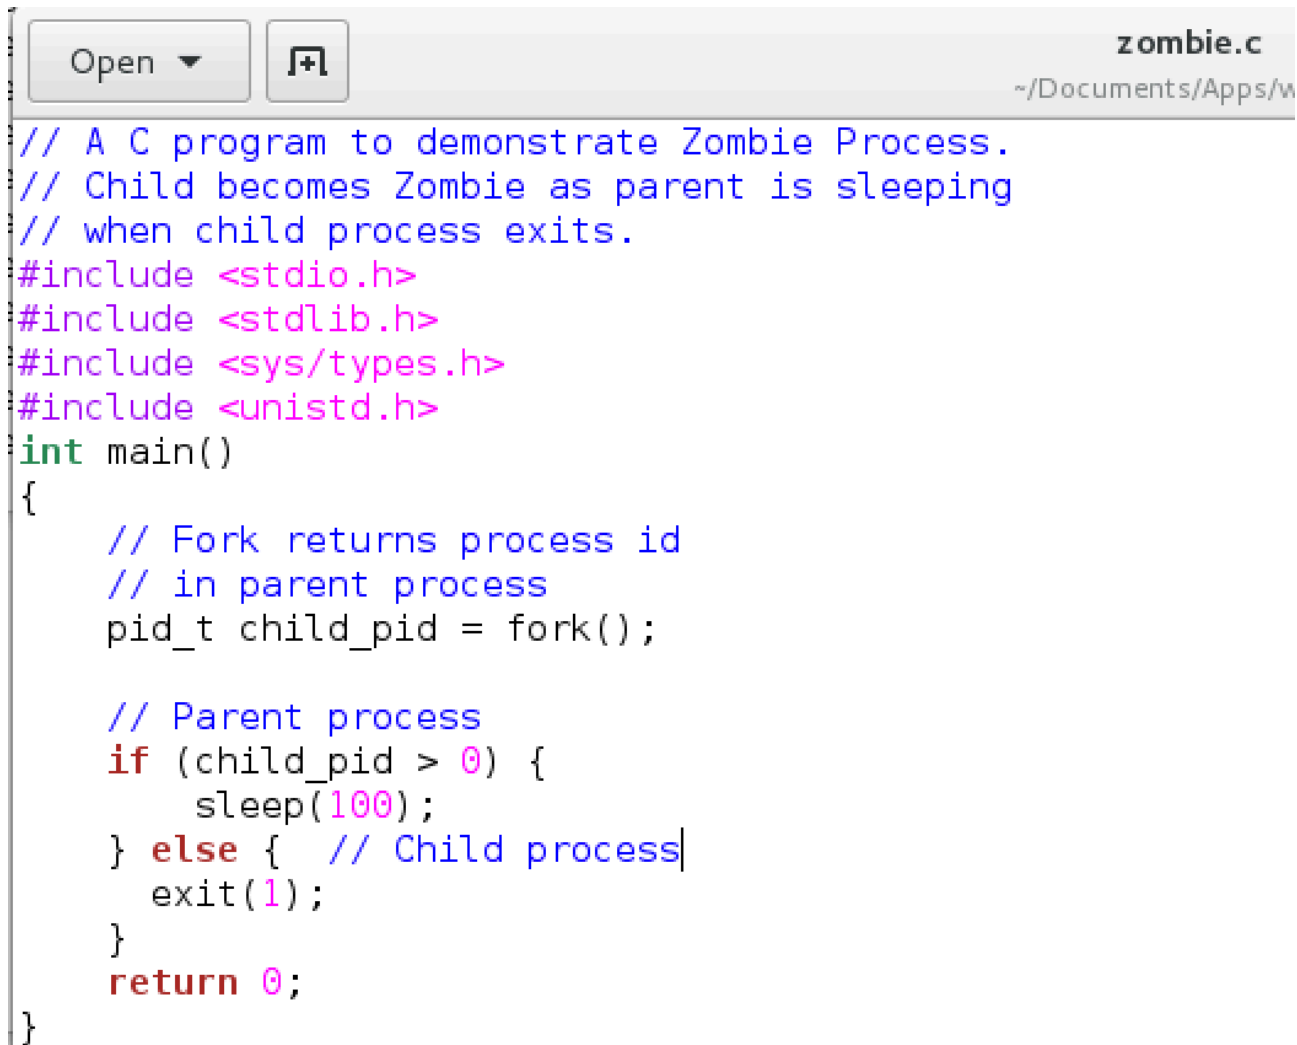
jmccarthy@debianJMC2017: ~/Documents/Apps/week4
File Edit View Search Terminal Help
1398  1498  1046  1046 ?         -1  S     1000   0:00 gnome-pty-helper
1398  1499  1499  1499 pts/0      1499 Ss+   1000   0:00 bash
1    1536  1046  1046 ?         -1  Sl    1000   0:00 /usr/lib/dconf/dconf-se
1145  2178  968   968 ?         -1  Sl    1000   7:17 firefox-esr
1398  2458  2458  2458 pts/1      5348 Ss    1000   0:00 bash
1    3996  1046  1046 ?         -1  Sl    1000   0:15 /usr/bin/gedit --gappli
1398  4122  4122  4122 pts/2      4122 Ss+   1000   0:00 bash
1    5343  5342  1499 pts/0      1499 S     1000   0:00 ./orphan
2458  5348  5348  2458 pts/1      5348 R+    1000   0:00 ps -xj
jmccarthy@debianJMC2017:~/Documents/Apps/week4$

```

# Zombie Processes

- When a process terminates, it isn't removed straight away from memory.
- The process status becomes `EXIT_ZOMBIE` and its parent is notified via `SIGCHLD` signal.
- The parent should execute the `wait()` system call to read the child processes exit status.
- Once the parent has processed this information from the child, the terminated process can be removed from memory.
- If the parent is unable to process this information the child process becomes a Zombie Process.

# Zombie Example



```
// A C program to demonstrate Zombie Process.
// Child becomes Zombie as parent is sleeping
// when child process exits.
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
int main()
{
    // Fork returns process id
    // in parent process
    pid_t child_pid = fork();

    // Parent process
    if (child_pid > 0) {
        sleep(100);
    } else { // Child process
        exit(1);
    }
    return 0;
}
```

# Zombie Example

```

jmccarthy@debianJMC2017: ~/Documents/Apps/week4
File Edit View Search Terminal Help
968 1233 968 968 ? -1 SNL 1000 0:00 /usr/lib/tracker/tracker
1 1238 1046 1046 ? -1 SL 1000 0:01 /usr/lib/tracker/tracker
1215 1239 1046 1046 ? -1 S 1000 0:00 /bin/cat
1 1265 1046 1046 ? -1 S 1000 0:00 /usr/lib/x86_64-linux-g
1 1332 1046 1046 ? -1 SL 1000 0:05 /usr/bin/nautilus --gap
1 1338 1046 1046 ? -1 SL 1000 0:00 /usr/lib/gvfs/gvfsd-tra
1 1347 1046 1046 ? -1 SL 1000 0:00 /usr/lib/gvfs/gvfsd-bur
1 1359 1046 1046 ? -1 SL 1000 0:00 /usr/lib/gvfs/gvfsd-met
1 1398 1046 1046 ? -1 SL 1000 0:21 /usr/lib/gnome-terminal
1398 1498 1046 1046 ? -1 S 1000 0:00 gnome-pty-helper
1398 1499 1499 1499 pts/0 5510 Ss 1000 0:00 bash
1 1536 1046 1046 ? -1 SL 1000 0:00 /usr/lib/dconf/dconf-se
1145 2178 968 968 ? -1 SL 1000 7:20 firefox-esr
1398 2458 2458 2458 pts/1 5513 Ss 1000 0:00 bash
1 3996 1046 1046 ? -1 SL 1000 0:18 /usr/bin/gedit --gappli
1398 4122 4122 4122 pts/2 4122 Ss+ 1000 0:00 bash
1499 5510 5510 1499 pts/0 5510 S+ 1000 0:00 ./zombie
5510 5511 5510 1499 pts/0 5510 Z+ 1000 0:00 [zombie] <defunct>
2458 5513 5513 2458 pts/1 5513 R+ 1000 0:00 ps -xj
jmccarthy@debianJMC2017:~/Documents/Apps/week4$

```

➤ Note the Process state codes (Z+)

➤ Defunct ("zombie") process, terminated but not reaped by its parent.

# Daemons

- A Daemon is a process.
- It runs in the background not under the control of a user.
- A daemon has a PID of 1. They are usually started when the system is booted and will only terminate on shutdown.
- Example: crond, ftpd, rlogind, mysqld, apache



# Steps to create a daemon

- Step 1: Create the orphan process
- Step 2: Elevate the orphan process to session leader, to loose controlling TTY
- Step 3: call `umask()` to set the file mode creation mask to 0  
This will allow the daemon to read and write files with the permissions/access required
- Step 4: Change the current working dir to root. This will eliminate any issues of running on a mounted drive, that potentially could be removed etc..
- Step 5: Close all open file descriptors

# Step 1: Create the orphan process

```

Open  [icon] step1.c
~/Documents/Apps/week4/ClassExample

// Orphan Example
// The child process is adopted by init process, when parent process dies.
#include<stdio.h>
#include <stdlib.h>

int main()
{
    // Implementation for Singleton Pattern if desired (Only one instance running)

    // Create a child process
    int pid = fork();

    if (pid > 0) {
        // if PID > 0 :: this is the parent
        // this process performs printf and finishes
        printf("Parent process");
        sleep(10); // uncomment to wait 10 seconds before process ends
        exit(EXIT_SUCCESS); // Kill the parent, needed to make orphan
    } else if (pid == 0) {
        // Step 1: Create the orphan process
        printf("Child process");

        // Orphan Logic goes here!!
        // Keep process running with infinite loop
        // When the parent finishes after 10 seconds,
        // the getppid() will return 1 as the parent (init process)
        while(1) {
            sleep(1);
            printf("child 1: my parent is: %i\n", getppid());
        }
    }

    return 0;
}

```

fork() to create an orphan.  
The orphan will not be a group leader.

The orphan will have the GPID of the parent process.

This is important, the next step will fail if the process is a Group Leader.

## Step 2: Elevate the orphan process to session leader

```
// Step 2: Elevate the orphan process to session leader, to loose controlling TTY
// This command runs the process in a new session
if (setsid() < 0) { exit(EXIT_FAILURE); }
```

- the setsid() system call is used to get the process to be the session group leader and process group leader.
- Call setsid() to get the process to become a process group leader and session group leader.
- The new session will not be associated with a controlling terminal

## Step 3: call umask() to set the file mode creation mask to 0

```
// Step 3: call umask() to set the file mode creation mask to 0
// This will allow the daemon to read and write files
// with the permissions/access required
umask(0);
```

- umask(0) so that we have the ability to read and write files etc. If this is not done correctly the daemon may not have the appropriate permissions to access files.

## Step 4: Change the current working dir to root.

```
// Step 4: Change the current working dir to root.  
// This will eliminate any issues of running on a mounted drive,  
// that potentially could be removed etc..  
if (chdir("/") < 0 ) { exit(EXIT_FAILURE); }
```

➤ `chdir("/")` to ensure that our process doesn't keep any directory in use.

# Step 5: Close all open file descriptors

```
// Step 5: Close all open file descriptors
/* Close all open file descriptors */
int x;
for (x = sysconf(_SC_OPEN_MAX); x>=0; x--)
{
    close (x);
}
```

- close() fds 0, 1, and 2.
- We are trying to release the IO connections inherited from the parent.
- use sysconf() to determine the limit \_SC\_OPEN\_MAX. This gives the max value of file descriptors (int value).
- Create a loop and close all open file descriptors.

# Process State Codes

## PROCESS STATE CODES

Here are the different values that the `s`, `stat` and `state` output specifiers (header "STAT" or "S") will display to describe the state of a process.

- D Uninterruptible sleep (usually IO)
- R Running or runnable (on run queue)
- S Interruptible sleep (waiting for an event to complete)
- T Stopped, either by a job control signal or because it is being traced.
- W paging (not valid since the 2.6.xx kernel)
- X dead (should never be seen)
- Z Defunct ("zombie") process, terminated but not reaped by its parent.

- PPID – Parent Process ID
- PID – Process ID
- PGID – Process Group ID
- SID – Session ID
- TTY – Controlling Terminal
- TPGID – Controlling tty process group ID
- STAT – Process Status Codes
- UID – Effective User Id



# Process Status Codes

## PROCESS STATE CODES

Here are the different values that the **s**, **stat** and **state** output specifiers (header "STAT" or "S") will display to describe the state of a process:

D	uninterruptible sleep (usually IO)
R	running or runnable (on run queue)
S	interruptible sleep (waiting for an event to complete)
T	stopped, either by a job control signal or because it is being traced
W	paging (not valid since the 2.6.xx kernel)
X	dead (should never be seen)
Z	defunct ("zombie") process, terminated but not reaped by its parent

# Error Logging

- How can a daemon deal with error messages?
- It doesn't have access to `standard_error`, as it doesn't have a controlling terminal.
- One option could be to write the error messages to a file, but this could get messy if there were loads of daemons writing to loads of files.
- A central solution would be best.

# Error Logging

- Most daemons use the syslog function to generate log messages.

SYSLOG(3)

Linux Programmer's Manual

SYSLOG(3)

## NAME

closelog, openlog, syslog, vsyslog - send messages to the system logger

## SYNOPSIS

```
#include <syslog.h>
```

```
void openlog(const char *ident, int option, int facility);
```

```
void syslog(int priority, const char *format, ...);
```

```
void closelog(void);
```

```
#include <stdarg.h>
```

```
void vsyslog(int priority, const char *format, va_list ap);
```

# Walkthrough Example – In Class Demo



**Run.....**

# Required Reading

➤ **How To Write a UNIX Daemon by Dave Lennert** (Hewlett-Packard Company)

➤ <http://cjh.polyplex.org/software/daemon.pdf>

➤ **Advanced Programming in the Unix Environment** – Second Edition – W. Richard Stevens and Stephen A. Rago.

➤ **Chapter 13.** Daemon Processes

# Questions

