



Systems Software

Week 5: IPC and Pipes



Overview

- ↗ Inter Process Communication (IPC)
- ↗ File Descriptors
- ↗ Pipes
- ↗ Pipes Example
- ↗ Named Pipes (FIFO)

- ↗ Inter Process Communication (IPC)
- ↗ Facilitating communication between different processes.
- ↗ IPC can be offered between:
 - ↗ Related processes
 - ↗ Unrelated processes

File Streams and File Descriptors

- If we need to perform file operations (IO) to a file, there are 2 options available:
 - Streams
 - File Descriptors
- Streams are represented as `File *` objects.
- File descriptors are represented as objects of type `int`.

File Streams

- A stream offers a high level interface that is layered on top of the file descriptors.
- There are more features and functionality that comes with the stream interface for file IO.

File Descriptors

- File descriptors offer a primitive low level interface for IO operations.
- A file descriptor can connect to a file, a device (terminal), or a pipe or socket to communicate with another process.
- File descriptors should be used for IO with devices, and for nonblocking IO operations.

fopen and fdopen

- The **fopen()** function opens the file whose name is the string pointed to by path and associates a stream with it.
- The **fdopen()** function associates a stream with the existing file descriptor, fd. The mode of the stream (one of the values "r", "r+", "w", "w+", "a", "a+") must be compatible with the mode of the file descriptor.

fgetc

- `fgetc()` reads the next character from stream and returns it as an unsigned char cast to an int, or EOF on end of file or error.


fstat

- `stat()` retrieve information about the file pointed to by `pathname`; the differences for `fstatat()` are described below.
- `fstat()` is identical to `stat()`, except that the file about which information is to be retrieved is specified by the file descriptor `fd`.

What are Pipes

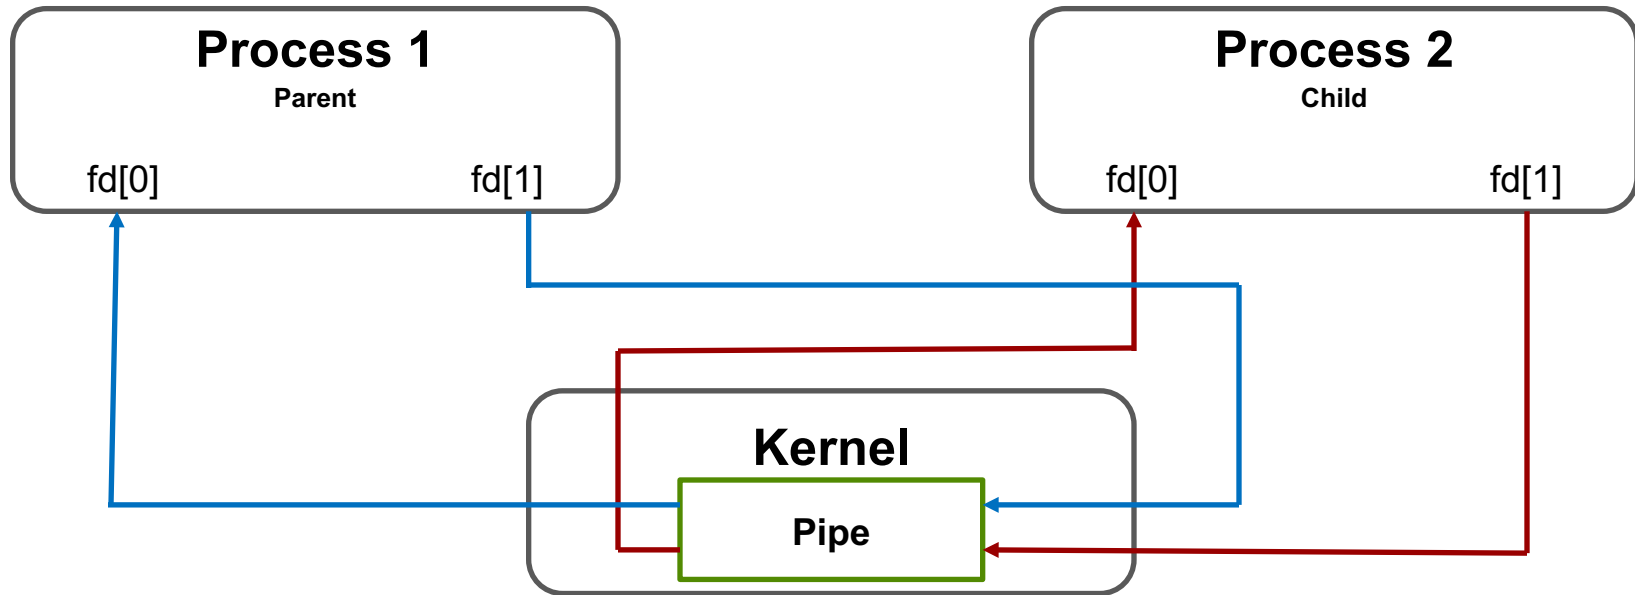
- A pipe allows processes to communicate with each other (Inter Process Communication - IPC)
- A process can write data to a pipe and this can be read by another process.
- The data passes to the pipe follows the FIFO algorithm.
- The pipe has no name.
- The pipe is created by a parent process and uses it to communicate with a child process.

Simple Pipe Architecture Example



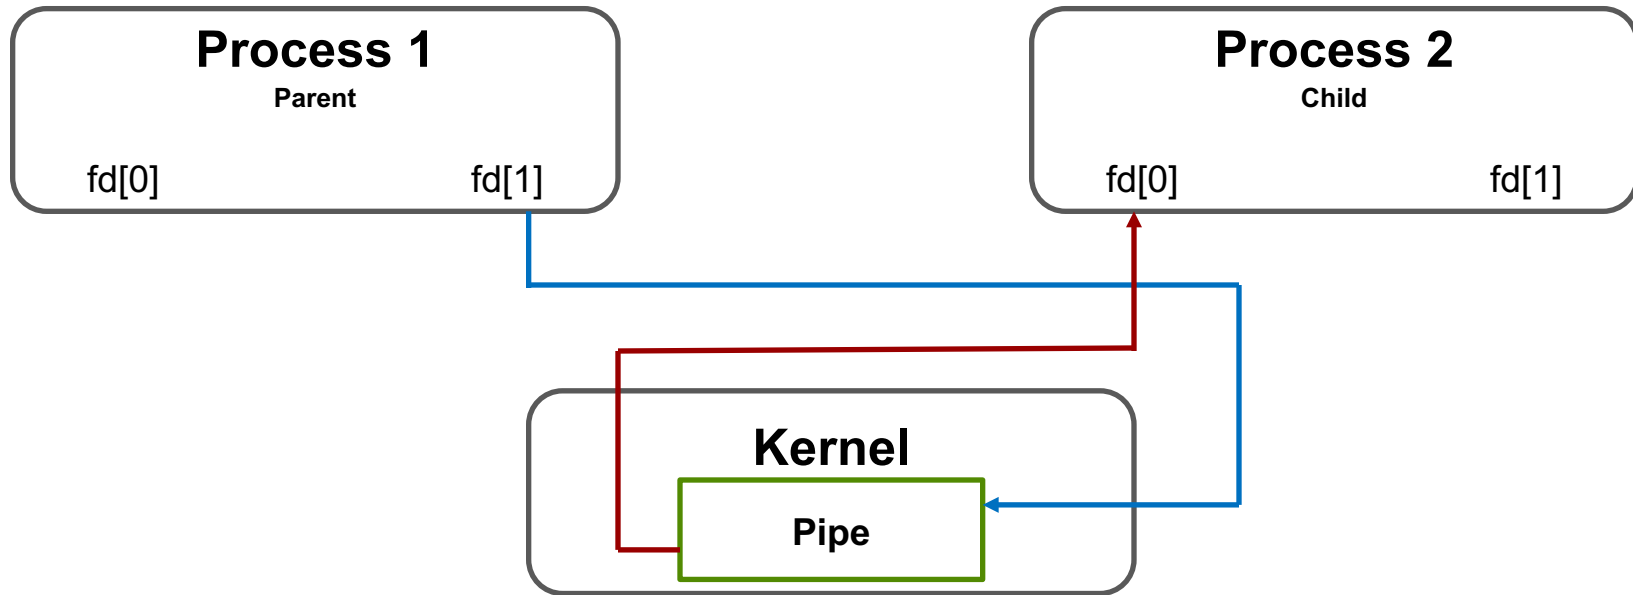
- Where fd = file descriptor
- fd[0] for reading
- fd[1] for writing
- The pipe is implemented using kernel memory.
- A pipe is a channel with two ends
- This example is a half-duplex pipe. Why??
- This example is of no real practical use, why would a single process need to communicate with itself?

Process forks



- After a fork we have to decide in which direction the data should flow
 - Parent to child
 - Child to parent

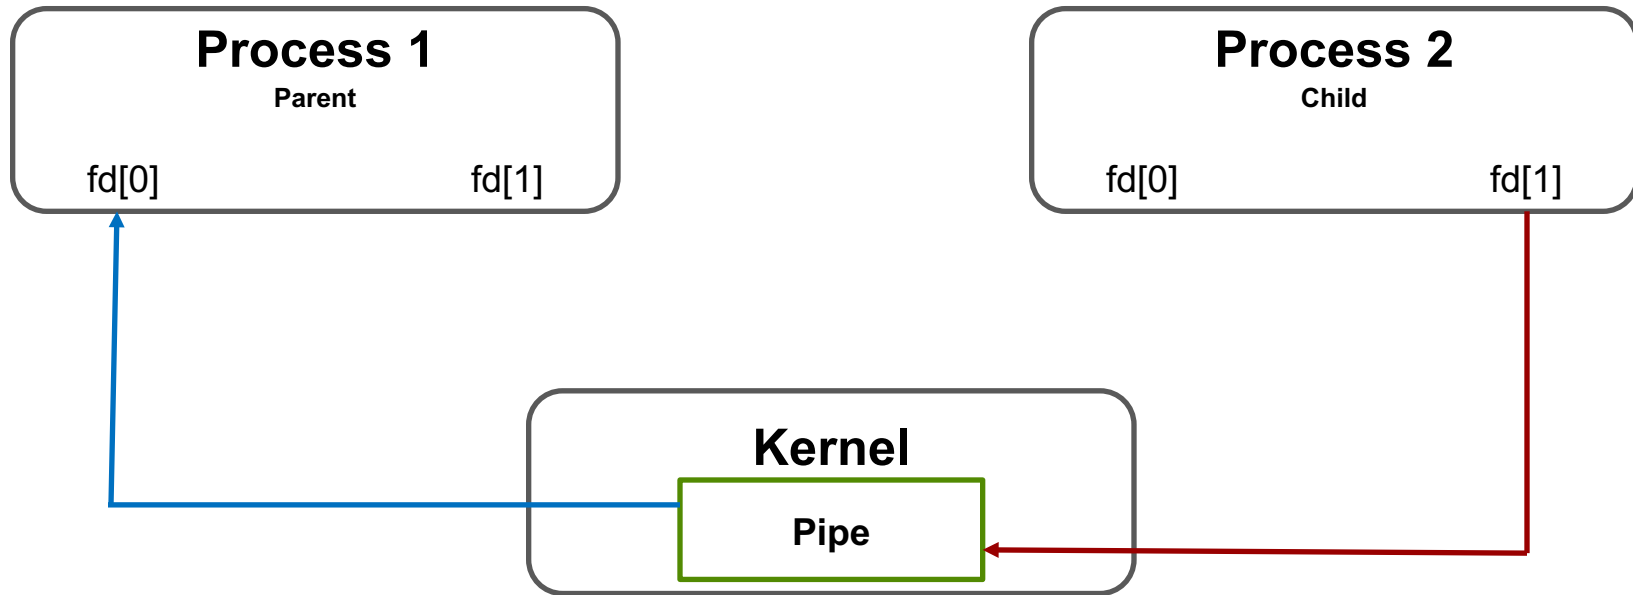
Parent to Child



➤ Parent must close `fd[0]`

➤ Child must close `fd[1]`

Child to Parent



➤ Parent must close fd[1]

➤ Child must close fd[0]

Rules for closing pipes::

- ↗ If we read from a pipe that was closed, the return should be 0
- ↗ If we write to a pipe with read closed on the other end, the SIGPIPE signal will be sent. This will need to be dealt with.

Creating a Pipe:: Step 1

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <stdlib.h>
#include <string.h>

int main(void)
{
```


Step 2:: Variables, pipe and fork

```
int      fd[2], nbytes;
pid_t    pid;
char     string[] = "Howya Friend!!!\n";
char     readbuffer[100];

// Create the pipe
pipe(fd);

// create a child process
pid = fork();
```

Step 3: Setup Child file Descriptors

```
if(pid == 0) // This is the child
{
    // Take no input, close fd[0]
    close(fd[0]);

    // Send output to parent via write
    write(fd[1], string, (strlen(string)+1));
    exit(0);
}
```

Step 4: Setup Child file Descriptors

```
else // This is the parent
{
    // Send no output, close fd[1]
    close(fd[1]);

    // Get input from the pipe via read
    nbytes = read(fd[0], readbuffer, sizeof(readbuffer));
    printf("Message from Child: %s", readbuffer);
}
```

Run

jmccarthy@debianJMC2017: ~/Documents/Apps/week5/pipes

x

File Edit View Search Terminal Help

```
jmccarthy@debianJMC2017:~/Documents/Apps/week5/pipes$ gcc -o pipe2 pipes2.c
jmccarthy@debianJMC2017:~/Documents/Apps/week5/pipes$ ./pipe2
Message from Child: Howya Friend!!!
jmccarthy@debianJMC2017:~/Documents/Apps/week5/pipes$
```

Pipe Example 2

- ↗ Create a C program to replicate the following:
 - ↗ **ps aux | grep login**
- ↗ How can pipes be used to solve this problem??

dup command

DUP(2)

Linux Programmer's Manual

DUP(2)

NAME

dup, dup2, dup3 - duplicate a file descriptor

SYNOPSIS

```
#include <unistd.h>
```

```
int dup(int oldfd);
```

```
int dup2(int oldfd, int newfd);
```

```
#define _GNU_SOURCE
```

```
#include <fcntl.h>
```

```
#include <unistd.h>
```

```
/* See feature_test_macros(7) */
```

```
/* Obtain 0_* constant definitions */
```

```
int dup3(int oldfd, int newfd, int flags);
```

dup command

dup2()

The **dup2()** system call performs the same task as **dup()**, but instead of using the lowest-numbered unused file descriptor, it uses the descriptor number specified in newfd. If the descriptor newfd was previously open, it is silently closed before being reused.

The steps of closing and reusing the file descriptor newfd are performed atomically. This is important, because trying to implement equivalent functionality using **close(2)** and **dup()** would be subject to race conditions, whereby newfd might be reused between the two steps. Such reuse could happen because the main program is interrupted by a signal handler that allocates a file descriptor, or because a parallel thread allocates a file descriptor.

Note the following points:

- * If oldfd is not a valid file descriptor, then the call fails, and newfd is not closed.
- * If oldfd is a valid file descriptor, and newfd has the same value as oldfd, then **dup2()** does nothing, and returns newfd.

dup command

`dup3()`

`dup3()` is the same as `dup2()`, except that:

- * The caller can force the close-on-exec flag to be set for the new file descriptor by specifying `O_CLOEXEC` in flags. See the description of the same flag in `open(2)` for reasons why this may be useful.
- * If oldfd equals newfd, then `dup3()` fails with the error `EINVAL`.

Pipes IPC Example 2

```
#include <stdlib.h>
#include <unistd.h>
```

```
void exec1();
void exec2();
```

```
int pid;
int pipefd[2];
```

- Init variables and functions.
- We have two functions that will use the exec command to perform different tasks in the process.
- Pid stores the process ID
- Pipefd array stores the file descriptors.

Pipes IPC Example 2

```
void main() {  
    // create pipe1  
    if (pipe(pipefd) == -1) {  
        perror("Error Init Pipe");  
        exit(1);  
    }  
}
```

- Start the program (main)
- Create a pipe and store the associated file descriptors in pipefd array

Pipes IPC Example 2

```
// fork (ps aux)
if ((pid = fork()) == -1) {
    perror("Error Init Fork 1");
    exit(1);
} else if (pid == 0) {
    exec1();
}
```

- Fork to create a child process
- If the pid is 0, we are dealing with the child process. Call the exec1 function to get the child process to swap to a different task.
- (exec1 function described on slide 27)

Pipes IPC Example 2

```
// fork (grep login)
if ((pid = fork()) == -1) {
    perror("Error Init Fork 1");
    exit(1);
} else if (pid == 0) {
    exec2();
}

close(pipefd[0]);
close(pipefd[1]);
} // close main
```

- Fork again to create a child process
- If the pid is 0, we are dealing with the child process. Call the exec2 function to get the child process to swap to a different task.
- (exec1 function described on slide 28)



Pipes IPC Example 2

```
void execl() {  
    // input from stdin - ok  
    // output to pipe  
    dup2(pipefd[1], 1);  
    // close fds  
    close(pipefd[0]);  
    close(pipefd[1]);  
    // exec  
    execlp("ps", "ps", "aux", NULL);  
    // exec didn't work, exit  
    perror("Error with ps aux");  
    exit(1);  
}
```

Pipes IPC Example 2

```
void exec2() {  
    // input from pipe  
    dup2(pipefd[0], 0);  
    // output to stdout - default  
  
    // close fds  
    close(pipefd[0]);  
    close(pipefd[1]);  
  
    // exec  
    execlp("grep", "grep", "login", NULL);  
    // exec didn't work, exit  
    perror("Error with grep login");  
    exit(1);  
}
```

Pipes – Making Life Easier!!

- The `popen()` function can be used to automate the creation of the pipe and the forking process.
- `fp = popen("ls *", "r");`
- `Popen` will set the file descriptors to facilitate the communication between parent and child or child and parent as required.

Pipes – Making Life Easier!!

popen.c x

*otherProcess.c x

```
#include <stdio.h>

int main() {
    FILE *fp;
    int status;
    char path[1024];

    fp = popen("ls *", "r");

    while (fgets(path, 1024, fp) != NULL)
        printf("%s", path);

    status = pclose(fp);
}
```

➤ Popen

1. Setup Pipes
2. Fork()
3. Child process runs `ls *` command and returns results to parent
4. Parent reads with `fgets` and displays with `printf`

Named Pipes

- A named pipe operates like a normal pipe
- The main difference is the named pipe exists as a special file
- The special file is a FIFO file, and doesn't contain any user info
- Processes of different ancestry can share data using named pipes.


FIFO file for a Named Pipe

- A named pipe operates like a file
- We can use the standead file IO system calls to operate a named pipe.
 - `int open(const char *pathname, int flags);`
 - `int read(int fd, void *buf, size_t count);`
 - `int write(int fd, const void *buf, size_t count);`
 - `int close(fd);`

FIFO Named Pipe using mkfifo

- The mkfifo function can be used to create a FIFO file.
- **mkfifo(fifoFile, 0666);**
- Use read and write operations to send information between processes.
- The processes don't need to be related (ancestors)
- Unlink can be called to remove the file.
- **unlink(fifoFile);**

FIFO Named Pipe Example



```
#include <fcntl.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int fd;
    char * fifoFile = "/tmp/fifoFile";

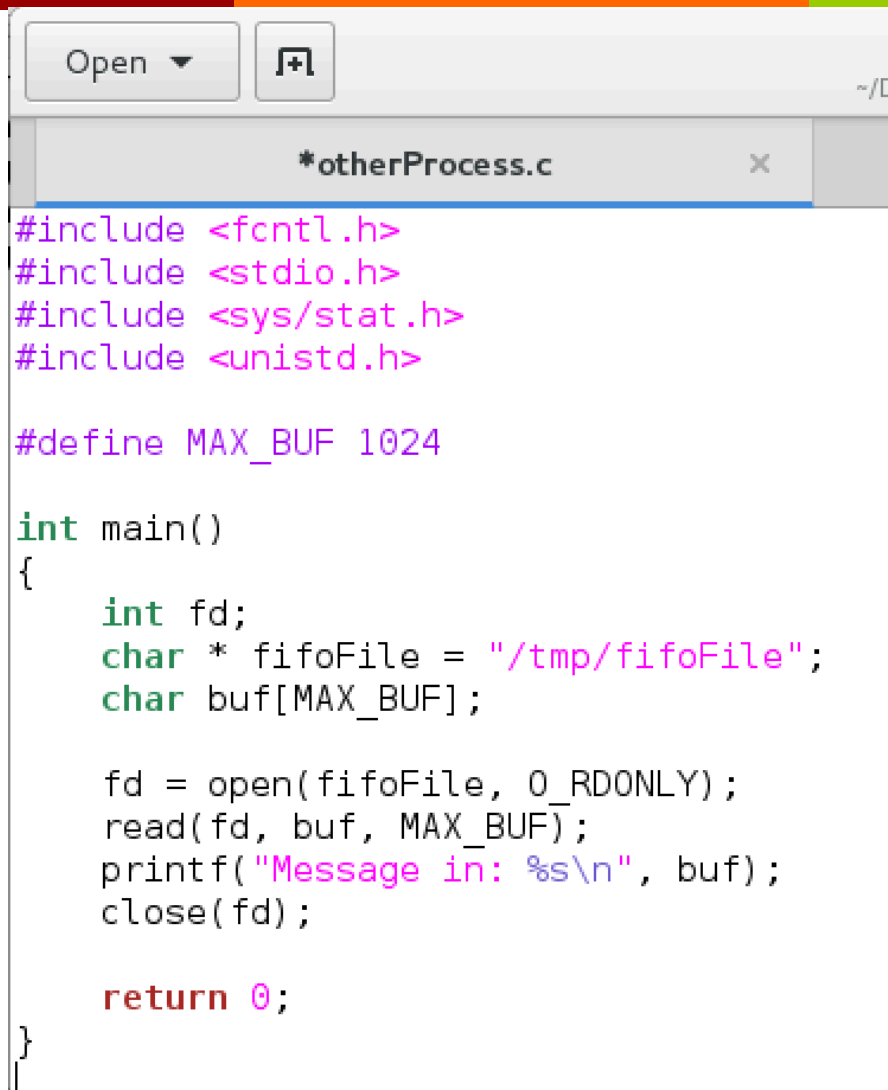
    /* create the FIFO (named pipe) */
    mkfifo(fifoFile, 0666);

    fd = open(fifoFile, O_WRONLY);
    write(fd, "The truth is out there!!", sizeof("The truth is out there!!"));
    close(fd);

    unlink(fifoFile);

    return 0;
}
```

FIFO Named Pipe Example



The image shows a code editor window with a title bar containing an "Open" button, a file icon, and a path indicator "~/.D". The editor has a tab labeled "*otherProcess.c" with a close button. The code is written in C and uses syntax highlighting: preprocessor directives are in magenta, keywords in green, and strings in magenta. The program defines a buffer size of 1024, opens a FIFO file for reading only, reads data into the buffer, prints it, and then returns 0.

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/stat.h>
#include <unistd.h>

#define MAX_BUF 1024

int main()
{
    int fd;
    char * fifoFile = "/tmp/fifoFile";
    char buf[MAX_BUF];

    fd = open(fifoFile, O_RDONLY);
    read(fd, buf, MAX_BUF);
    printf("Message in: %s\n", buf);
    close(fd);

    return 0;
}
```

Run Example

```
jmccarthy@debianJMC2017: ~/Documents/A
File Edit View Search Terminal Help
$ ./mainProcess
$
```

```
jmccarthy@debianJMC2017: ~/Docu
File Edit View Search Terminal Help
$ ./otherProcess
Message in: The truth is out there!!
$
```

Questions

